



The University of Manchester

Compatibility and Isolation in Specialised Operating Systems

Pierre Olivier

The University of Manchester

pierre.olivier@manchester.ac.uk

17/05/2023

Introducing Unikernels

HermiTux: a Unikernel Binary-Compatible with Linux

FlexOS: an Operating System for Flexible Isolation

Introducing Unikernels

HermiTux: a Unikernel Binary-Compatible with Linux

FlexOS: an Operating System for Flexible Isolation

Introductory example:
my website in the cloud

Introducing Unikernels

Presentation

Full-fledged Virtual Machine



Cloud provider:




amazon
web services™

Introducing Unikernels

Presentation


Full-fledged Virtual Machine



Introducing Unikernels


Presentation

Full-fledged Virtual Machine




Introducing Unikernels

Presentation



Introducing Unikernels


Presentation



Introducing Unikernels

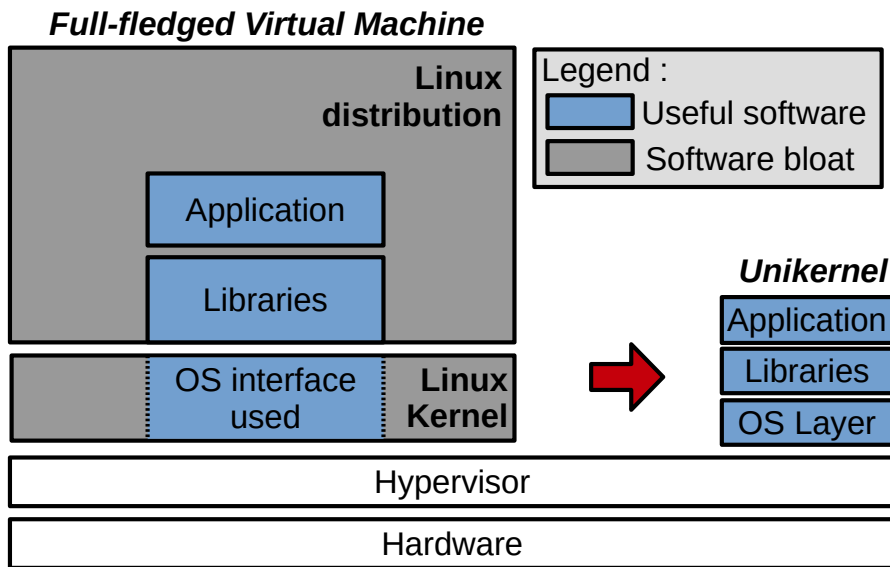
Presentation

Full-fledged Virtual Machine



Introducing Unikernels

Presentation



Introducing Unikernels

Presentation (2)

Unikernel: application + dependencies + thin OS compiled as a static binary running on top of a hypervisor ¹

¹Madhavapeddy et al., "Unikernels: Library Operating Systems for the Cloud", ASPLOS'13

²Zhang et al., "KylinX: A Dynamic Library Operating System for Simplified and Efficient Cloud Virtualization, ATC'18

Introducing Unikernels

Presentation (2)

Unikernel: application + dependencies + thin OS compiled as a static binary running on top of a hypervisor ¹

- **Single-purpose:** run 1 application
 - Want to run multiple applications? Run multiple unikernels
- **Single-process**
 - Want to run a multi-process application? Run multiple unikernels ²
 - However, SMP (multicores) and multithreading are supported
- **Single-binary and single-address space** for application + kernel
 - No user/kernel protection needed

¹Madhavapeddy et al., "Unikernels: Library Operating Systems for the Cloud", ASPLOS'13

²Zhang et al., "KylinX: A Dynamic Library Operating System for Simplified and Efficient Cloud Virtualization, ATC'18

- **Lightweight virtualisation**

- Contain and run only what is absolutely necessary for the application
- Security advantage: small attack surface
- Cost advantage: memory/disk footprint reduction
- As VMs, considered as a *secure alternative to containers*

- **Lightweight virtualisation**

- Contain and run only what is absolutely necessary for the application
- Security advantage: small attack surface
- Cost advantage: memory/disk footprint reduction
- As VMs, considered as a *secure alternative to containers*

- **Per-application tailored kernel**

- LibOS/Exokernel model, specialised subsystems
- The kernel itself contains only what is needed

- **Lightweight virtualisation**

- Contain and run only what is absolutely necessary for the application
- Security advantage: small attack surface
- Cost advantage: memory/disk footprint reduction
- As VMs, considered as a *secure alternative to containers*

- **Per-application tailored kernel**

- LibOS/Exokernel model, specialised subsystems
- The kernel itself contains only what is needed

- **Reduced OS noise, increased performance**

- Low system call latency: app + kernel in ring 0, system calls are function calls
- Sub-second boot time

Introducing Unikernels

Unikernels Models

Unikernels can be classified based on the language(s) in which the applications they support can be written in (non-exhaustive list):

- **Pure memory safe languages** (OCamL, Erlang, Haskell): MirageOS¹, LING², HaLVM³
- **C/C++, semi-POSIX API**: HermitCore⁴, OSv⁵, Rumprun⁶, Lupine⁷, UKL⁸
- **Rust/Go**: RustyHermit⁹, Clive¹⁰
- **Multi-languages**: HermitTux¹¹, Unikraft¹²

¹ <https://mirage.io/>, ² <https://github.com/cloudozer/ling> ³ <https://github.com/GaloisInc/HaLVM> ⁴ <https://hermitcore.org/> ⁵ <http://osv.io/> ⁶ <https://github.com/rumpkernel/rumprun> ⁷ <https://github.com/hckuo/Lupine-Linux> ⁸ <https://github.com/unikernelLinux/ukl> ⁹ <https://github.com/hermitcore/libhermit-rs> ¹⁰ <http://lsub.org/ls/clive.html> ¹¹ <https://ssrg-vt.github.io/hermitux/> ¹² <https://unikraft.org/>

Introducing Unikernels

Summary

Benefits Summary

- Lower costs by being lightweight
- Increased security with low attack surface and high isolation
- Increased performance with low OS disturbance and fast boot times

Introducing Unikernels

Summary

Benefits Summary

- Lower costs by being lightweight
- Increased security with low attack surface and high isolation
- Increased performance with low OS disturbance and fast boot times

Application domains

- Cloud applications: servers, micro-services, SaaS, FaaS/serverless
- Embedded virtualisation, Edge computing, IoT
- Network Function Virtualisation, HPC, VM introspection, malware analysis, secure Desktop applications
- etc.

Introducing Unikernels

HermiTux: a Unikernel Binary-Compatible with Linux

FlexOS: an Operating System for Flexible Isolation

HermiTux: a Unikernel Binary-Compatible with Linux

The Issue

- Unikernels have plenty of benefits to bring
- Unikernels have plenty of application domains
- They are very popular in academia ...
 - Mirage (ASPLOS'13), LightVM (SOSP'17), and many others
- **...but why (nearly) nobody uses them in the industry?**

HermiTux: a Unikernel Binary-Compatible with Linux


The Issue

- Unikernels have plenty of benefits to bring
- Unikernels have plenty of application domains
- They are very popular in academia ...
 - Mirage (ASPLOS'13), LightVM (SOSP'17), and many others
- **...but why (nearly) nobody uses them in the industry?**

Because it is hard to port existing applications!


HermiTux: a Unikernel Binary-Compatible with Linux

Why Porting to Unikernels is Hard



HermiTux: a Unikernel Binary-Compatible with Linux


Why Porting to Unikernels is Hard



- Proprietary software → source code not available

HermiTux: a Unikernel Binary-Compatible with Linux


Why Porting to Unikernels is Hard



- Proprietary software → source code not available
- Incompatible language

HermiTux: a Unikernel Binary-Compatible with Linux

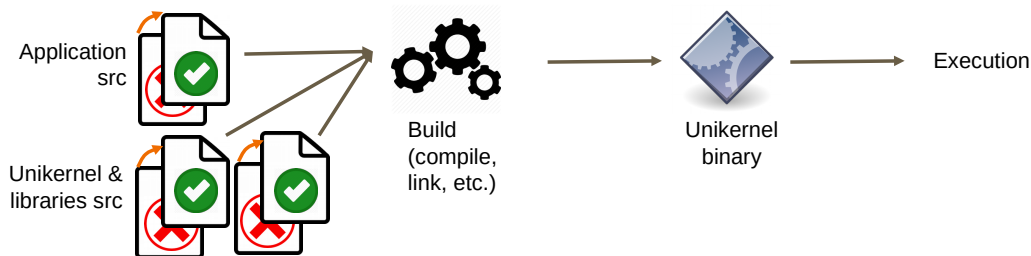
Why Porting to Unikernels is Hard



- Proprietary software → source code not available
- Incompatible language
- Unsupported features

HermiTux: a Unikernel Binary-Compatible with Linux


Why Porting to Unikernels is Hard



- Proprietary software → source code not available
- Incompatible language
- Unsupported features
- Porting is hard, needs knowledge about both application and unikernel

HermiTux: a Unikernel Binary-Compatible with Linux

Why Porting to Unikernels is Hard



- Proprietary software → source code not available
- Incompatible language
- Unsupported features
- Porting is hard, needs knowledge about both application and unikernel
- Complex build toolchains

HermiTux: a Unikernel Binary-Compatible with Linux

Why Porting to Unikernels is Hard

- Proprietary software → source code not available
- Incompatible language
- Unsupported features
- Porting is hard, needs knowledge about both application and unikernel
- Complex build toolchains

Our Solution: HermiTux

- A unikernel *binary-compatible* with Linux applications
 - I.e. a unikernel that can run application that have been compiled for Linux
 - For the x86-64 and aarch64 (ARM64) architectures



HermiTux: a Unikernel Binary-Compatible with Linux

Binary Compatibility with Linux Applications

HermiTux's objective: running as unikernels executables that have been compiled for a popular OS: Linux

HermiTux: a Unikernel Binary-Compatible with Linux

Binary Compatibility with Linux Applications

HermiTux's objective: running as unikernels executables that have been compiled for a popular OS: Linux

- **Consider unmodified binaries** built with various compilers, from various source languages, potentially stripped/obfuscated
- **Don't assume access to applications' source code**
- **Consider both static and dynamic binaries**
- **Maintain unikernels principles and benefits**

HermiTux: a Unikernel Binary-Compatible with Linux

Binary Compatibility with Linux Applications

HermiTux's objective: running as unikernels executables that have been compiled for a popular OS: Linux

- **Consider unmodified binaries** built with various compilers, from various source languages, potentially stripped/obfuscated
 - **Don't assume access to applications' source code**
 - **Consider both static and dynamic binaries**
 - **Maintain unikernels principles and benefits**
-
- Binary compatibility is achieved by developing a unikernel **conforming to the Linux Application Binary Interface (ABI)**
 - Convention, sets of rules a program needs to follow to execute on top of the Linux kernel
 - We need to implement the OS side of these rules in our unikernel

HermiTux: a Unikernel Binary-Compatible with Linux

The Linux ABI

It is partially architecture (ISA) specific, and is composed of two main sets of rules:

HermiTux: a Unikernel Binary-Compatible with Linux

The Linux ABI

It is partially architecture (ISA) specific, and is composed of two main sets of rules:

- **Load-time conventions**

- Application binary format supported: ELF
- What part of the address space is accessible for the application (currently lower half of the 48 bits virtual address space)
- Layout of the application's stack & registers' content at the program entry point

HermiTux: a Unikernel Binary-Compatible with Linux

The Linux ABI

It is partially architecture (ISA) specific, and is composed of two main sets of rules:

- **Load-time conventions**

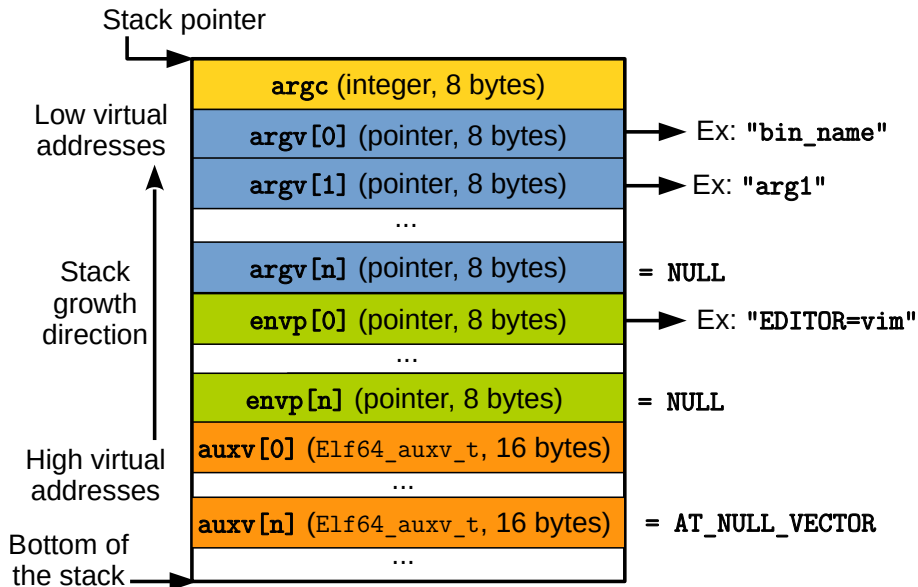
- Application binary format supported: ELF
- What part of the address space is accessible for the application (currently lower half of the 48 bits virtual address space)
- Layout of the application's stack & registers' content at the program entry point

- **Runtime conventions**

- How to request services from the OS: system calls
- Additional communication channels with the kernel through virtual filesystems (`/sys`, `/proc`, etc.) and shared memory areas (`vDSO/vsyscall`)

HermiTux: a Unikernel Binary-Compatible with Linux

Load-Time Convention: Stack at Entry Point



HermiTux: a Unikernel Binary-Compatible with Linux

Runtime Convention: System Calls

- Describes syscall invocation, parameters and return value passing.
- It is ISA-specific
- E.g. for x86-64:
 - Syscall identifier in `%rax`
 - Parameters are passed in orders in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
 - Return value in `%rax`
 - Invocation with the `syscall` instruction

HermiTux: a Unikernel Binary-Compatible with Linux

Overview

- Custom KVM-based hypervisor


The diagram consists of two horizontal rectangular boxes. The top box is light blue and contains the text 'Hypervisor: uHyve'. A black arrow points downwards from the center of this box to the center of the bottom box. The bottom box is light green and contains the text 'Host: Linux kernel' followed by 'KVM' in an italicized font.

Hypervisor: uHyve

Host: Linux kernel *KVM*

HermiTux: a Unikernel Binary-Compatible with Linux

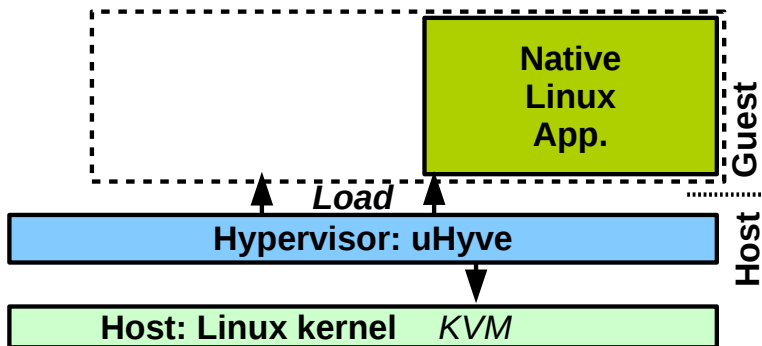
Overview



- Custom KVM-based hypervisor

HermiTux: a Unikernel Binary-Compatible with Linux

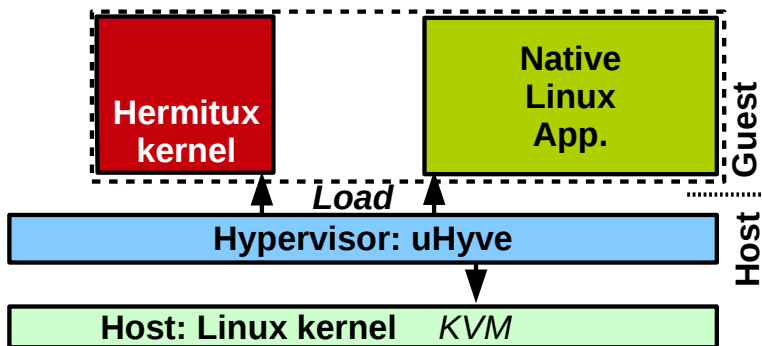
Overview



- Custom KVM-based hypervisor

HermiTux: a Unikernel Binary-Compatible with Linux

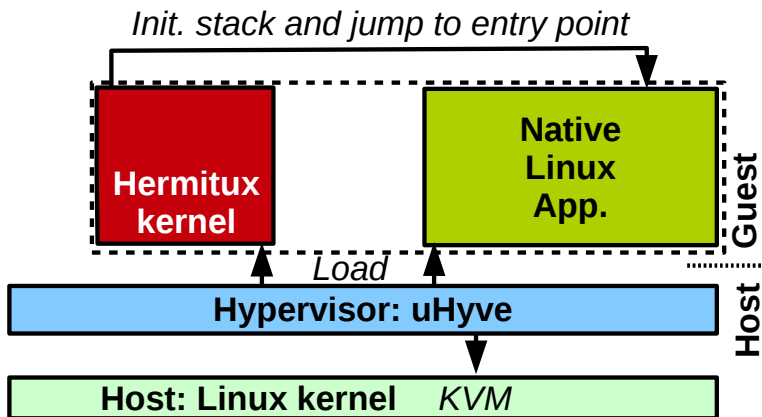
Overview



- Custom KVM-based hypervisor
- VMM loads app and kernel ELF binaries

HermiTux: a Unikernel Binary-Compatible with Linux


Overview



- Custom KVM-based hypervisor
- VMM loads app and kernel ELF binaries
- Follows load time ABI conventions

HermiTux: a Unikernel Binary-Compatible with Linux


Overview



- Custom KVM-based hypervisor
- VMM loads app and kernel ELF binaries
- Follows load time ABI conventions
- And runtime conventions for syscalls

HermiTux: a Unikernel Binary-Compatible with Linux

Overview



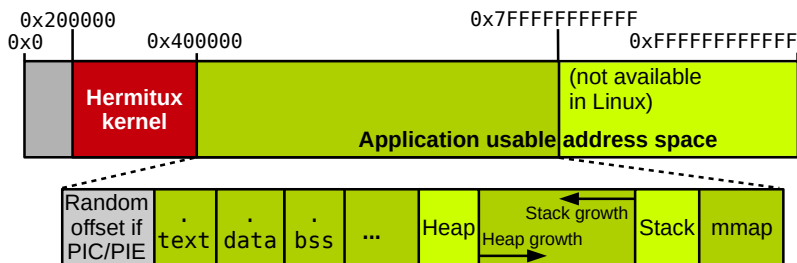
- Custom KVM-based hypervisor
- VMM loads app and kernel ELF binaries
- Follows load time ABI conventions
- And runtime conventions for syscalls

HermiTux: a Unikernel Binary-Compatible with Linux

Features: System Call Support

HermiTux supports 97 syscalls, 81 developed on top of the original HermitCore kernel. They offer support for:

- Statically and dynamically compiled binaries
- Multithreading (TLS, clone, synchronisation with futex)
- Signals
- Scheduling priorities
- Highly randomised memory mappings (entropy: 34 bits)



HermiTux: a Unikernel Binary-Compatible with Linux

Features: System Call Support (2)

Category	System calls supported by HermiTux
Filesystem	access, chdir, close, creat, faccessat, fcntl, fdatsync, fstat, fsync, getcwd, getdents, getdents64, lseek, lstat, mkdir, mkdirat, newfstatat, open, openat, pwrite64, read, readlinkat, readlink, readv, rmdir, stat, sync, syncfs, unlink, unlinkat, write, writev
Memory management	brk, sbrk, madvise, mincore, mmap, mprotect, mremap, munmap
Process management	clone*, exit, exit_group, getpid/ppid/gid/egid/euid/tid/uid, setsid, getrlimit, prlimit64, setrlimit, umask
Networking	accept, bind, connect, gethostname, getpeername, getsockname, getsockopt, listen, recvfrom, select, sendto, sendfile, sethostname, setsockopt, socket
Signals & Synchronization	futex, get_robust_list, kill, set_tid_address, rt_sigaction, sigaltstack, signal, tkill, tkill
Scheduling	getprio, setprio, sched_getaffinity, sched_setaffinity, sched_yield
Time management	clock_getres, clock_gettime, gettimeofday, nanosleep, time
Miscellaneous	arch_prctl, ioctl, shutdown, sysinfo, uname

*threads only

Unikernel Benefits in Hermitux

The Issue

- Until now we described a small operating system that can run applications and bring these unikernels benefits
 - Security and isolation
 - Low memory/disk footprint
 - Fast boot time

What about the other unikernel benefits?

- Fast system calls (function calls)
- Kernel modularity

Unikernel Benefits in Hermitux

Fast Syscalls for Dynamic Binaries with Libc Substitution


- Hermitux's syscall handler is invoked by the `syscall` instruction
 - Reintroduce high latency for system calls due to the exception

```
40127e:  b8 07 00 00 00      mov     $0x7,%eax (poll)
401283:  4c 89 c7             mov     %r8,%rdi
401286:  48 89 ca             mov     %rcx,%rdx
401289:  0f 05                syscall
```


Unikernel Benefits in HermiTux

Fast Syscalls for Dynamic Binaries with Libc Substitution

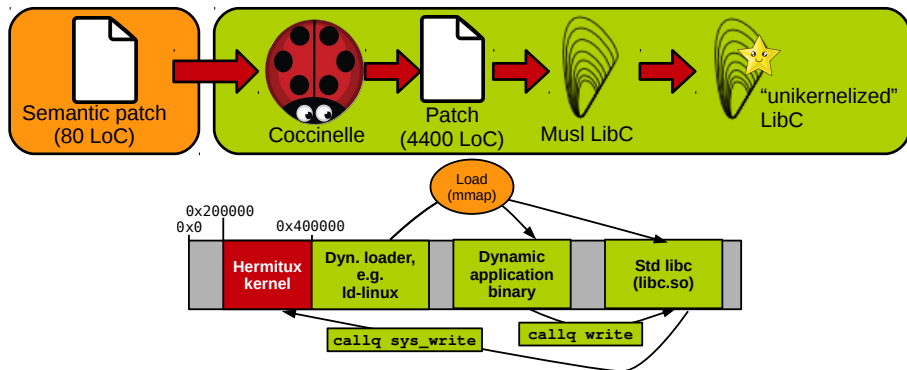
- For **dynamically compiled programs**:
 - At runtime **load a unikernel-aware Libc**
 - Custom Libc has system calls implemented as (fast) function calls directly into the kernel
 - Fork of Musl Libc, automatically transformed using Coccinelle



Unikernel Benefits in Hermitux

Fast Syscalls for Dynamic Binaries with Libc Substitution

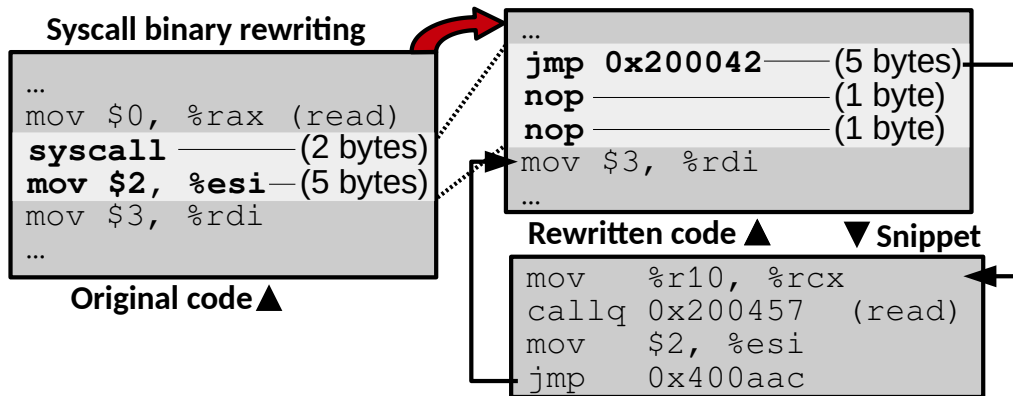
- For **dynamically compiled programs**:
 - At runtime **load a unikernel-aware Libc**
 - Custom Libc has system calls implemented as (fast) function calls directly into the kernel
 - Fork of Musl Libc, automatically transformed using Coccinelle



Unikernel Benefits in Hermitux

Fast Syscalls with Binary Rewriting

- What about **static binaries**?
- (Statically) **binary-rewrite syscall instructions to direct jumps to the syscall implementation**
 - Problem: syscall is 2 bytes long and any call/jmp instruction will be larger



Unikernel Benefits in Hermitux

Fast Syscalls with Binary Rewriting

- What about **static binaries**?
- (Statically) **binary-rewrite syscall** instructions to direct jumps to the syscall implementation
 - Problem: syscall is 2 bytes long and any call/jmp instruction will be larger

Syscall binary rewriting

```
...  
mov $0, %rax (read)  
syscall _____ (2 bytes)  
mov $2, %esi _____ (5 bytes)  
mov $3, %rdi  
...
```

Original code ▲

Different/smarter solutions appeared since then: X-Container (ASPLOS'19), NullPoline (ATC'23)

```
...  
jmp 0x200042 _____ (5 bytes)  
nop _____ (1 byte)  
nop _____ (1 byte)  
mov $3, %rdi  
...
```


Rewritten code ▲

▼ Snippet

```
mov %r10, %rcx  
callq 0x200457 (read)  
mov $2, %esi  
jmp 0x400aac
```

Unikernel Benefits in HermitTux

Fast Syscalls with Binary Rewriting



Unikernel Benefits in Hermitux

System-call-based Modularity

- System-call based modularity
 - Compile a kernel with support for only the necessary system calls
 - How to identify syscall needed without access to the sources?
 - **Use binary analysis** to find out what is the value in %rax (x86-64) or %x8 (aarch64) for each syscall invocation

	00 00 00			
00401bca	41 56	PUSH	R14	
00401bcc	41 55	PUSH	R13	
00401bce	41 bd 14	MOV	R13D, 0x14	
	00 00 00			
00401bd4	41 54	PUSH	R12	
00401bd6	49 89 d4	MOV	R12, RDX	
00401bd9	55	PUSH	RBP	
00401bda	53	PUSH	RBX	
00401bdb	48 89 fb	MOV	RBX, RDI	
00401bde	48 83 ec 28	SUB	RSP, 0x28	
00401be2	48 8b 47 38	MOV	RAX, qword ptr [RDI + 0x38]	
00401be6	4c 8b 77 28	MOV	R14, qword ptr [RDI + 0x28]	
00401bea	48 89 74	MOV	qword ptr [RSP + local_48], RSI	
	24 10			
00401bef	48 89 e5	MOV	RBX, RSP	
00401bf2	48 89 04 24	MOV	qword ptr [RSP] => local_58, RAX	
00401bf6	49 29 c6	SUB	R14, RAX	
00401bf9	48 89 54	MOV	qword ptr [RSP + local_40], RDX	
	24 18			
00401bfe	4c 89 74	MOV	qword ptr [RSP + local_50], R14	
	24 08			
00401c03	49 01 d6	ADD	R14, RDX	
		LAB_00401c06		
00401c06	48 63 7b 78	MOVSSD	RDI, dword ptr [RBX + 0x78]	
00401c0a	49 63 d7	MOVSSD	RDX, R15D	
00401c0d	4c 89 e8	MOV	RAX, R13	
00401c10	48 89 ee	MOV	RSI, RBP	
00401c13	0f 05	SYSCALL		
00401c15	48 89 c7	MOV	RDI, RAX	

Syscall id 0x14 → writev

XREF[1]: 00401c85(j)

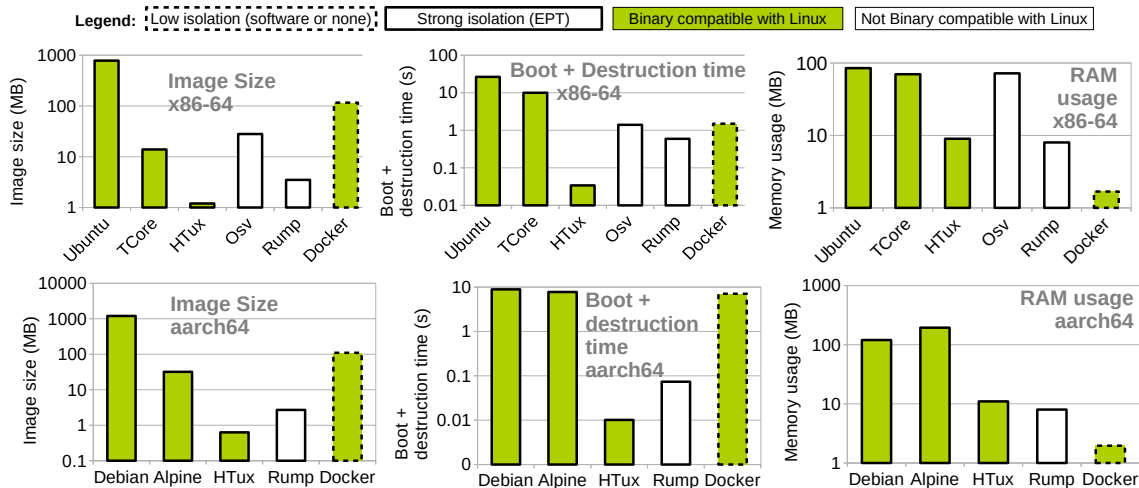
Unikernel Benefits in Hermitux

System-call-based Modularity

Program	Number of system calls	x86-64 kernel .text size reduction	aarch64 kernel .text size reduction
Minimal	4	21.22 %	29.26 %
Hello world	9	19.91 %	27.42 %
PARSEC Blackscholes	15	17.68 %	24.50 %
Postmark	27	16.02 %	22.55%
Sqlite	33	11.34 %	16.44%
Full syscalls support	97	-	-

HermiTux: a Unikernel Binary-Compatible with Linux


Evaluation



- Image 650x smaller, boot time 780x faster, RAM usage 9x lower than a Linux VM!


HermiTux: a Unikernel Binary-Compatible with Linux

Evaluation (2)





HermiTux: a Unikernel Binary-Compatible with Linux

Evaluation (3)




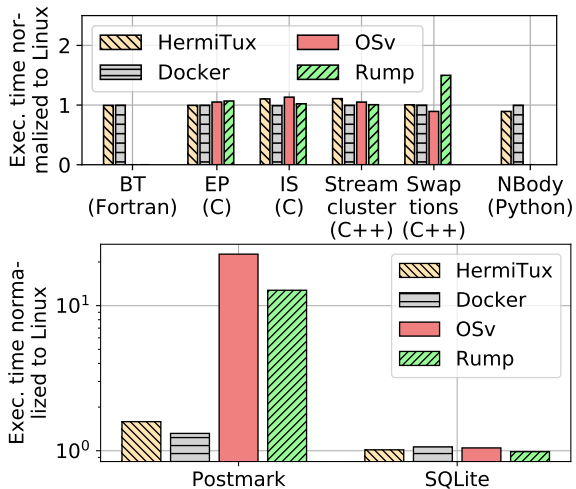
HermiTux: a Unikernel Binary-Compatible with Linux

Evaluation (3)







HermiTux: a Unikernel Binary-Compatible with Linux

Evaluation (3)



HermiTux: a Unikernel Binary-Compatible with Linux

Evaluation (3)



Demo

Introducing Unikernels

HermiTux: a Unikernel Binary-Compatible with Linux

FlexOS: an Operating System for Flexible Isolation

FlexOS: an Operating System for Flexible Isolation


Motivation

- OS security/isolation strategies are fixed at design time
 - Isolation granularity, mechanisms used, data sharing strategies, etc.

FlexOS: an Operating System for Flexible Isolation

Motivation


- OS security/isolation strategies are fixed at design time
 - Isolation granularity, mechanisms used, data sharing strategies, etc.



FlexOS: an Operating System for Flexible Isolation

Motivation


- OS security/isolation strategies are fixed at design time
 - Isolation granularity, mechanisms used, data sharing strategies, etc.
- Not really suitable in modern scenarios
 - Applications have heterogeneous needs in terms of security/performance
 - Application made of multiple components with various levels of trust
 - Machines support various isolation mechanisms, with new technologies underway
 - Hardware protection breaks (e.g. Meltdown)



FlexOS: an Operating System for Flexible Isolation

Motivation

- OS security/isolation strategies are fixed at design time
 - Isolation granularity, mechanisms used, data sharing strategies, etc.
- Not really suitable in modern scenarios
 - Applications have heterogeneous needs in terms of security/performance
 - Application made of multiple components with various levels of trust
 - Machines support various isolation mechanisms, with new technologies underway
 - Hardware protection breaks (e.g. Meltdown)



FlexOS

- LibOS that can specialise for **security**
- Security/isolation strategy can be instantiated at build time

FlexOS: an Operating System for Flexible Isolation


Flexible Builds

- Flexos decouple from the OS design important security/isolation decisions and allow selecting at build time:

FlexOS: an Operating System for Flexible Isolation

Flexible Builds


- Flexos decouple from the OS design important security/isolation decisions and allow selecting at build time:
 - Compartmentalisation strategy and granularity**, as well as various **data sharing strategies** for communicating compartments



FlexOS: an Operating System for Flexible Isolation

Flexible Builds


- Flexos decouple from the OS design important security/isolation decisions and allow selecting at build time:
 - Compartmentalisation strategy and granularity**, as well as various **data sharing strategies** for communicating compartments
 - Different levels of per-compartment **software hardening** (ASan, etc.)



FlexOS: an Operating System for Flexible Isolation

Flexible Builds

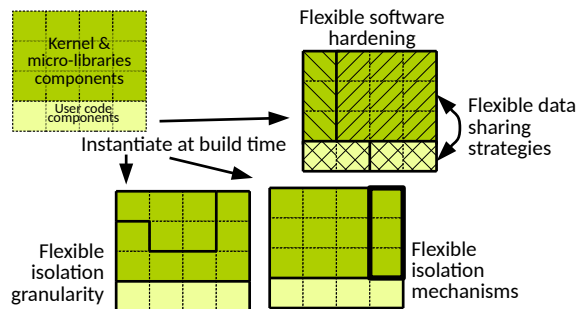
- Flexos decouple from the OS design important security/isolation decisions and allow selecting at build time:
 - Compartmentalisation strategy and granularity**, as well as various **data sharing strategies** for communicating compartments
 - Different levels of per-compartment **software hardening** (ASan, etc.)
 - Various **isolation mechanisms** to enforce the compartmentalisation: MPK, EPT



FlexOS: an Operating System for Flexible Isolation

Flexible Builds

- Flexos decouple from the OS design important security/isolation decisions and allow selecting at build time:
 - Compartmentalisation strategy and granularity**, as well as various **data sharing strategies** for communicating compartments
 - Different levels of per-compartment **software hardening** (ASan, etc.)
 - Various **isolation mechanisms** to enforce the compartmentalisation: MPK, EPT



How to enable flexible builds?

FlexOS: an Operating System for Flexible Isolation


Enabling Flexible Builds

- LibOS **compartmentalised to the finest level:** components
- Compartmentalisation API: annotations in the code marking **components interfaces** and **shared data**

FlexOS: an Operating System for Flexible Isolation

Enabling Flexible Builds

- LibOS compartmentalised to the finest level: components
- Compartmentalisation API: annotations in the code marking **components interfaces** and **shared data**
- **Code transformations** at build time merging/separating components into compartments and instantiating gates



FlexOS: an Operating System for Flexible Isolation

Enabling Flexible Builds: Porting API

```
int rc, connfd;  
char buf[512];  
/* ... */  
rc = recv(connfd, buf, 512, 0);
```

FlexOS: an Operating System for Flexible Isolation

Enabling Flexible Builds: Porting API

```
int rc, connfd;  
char buf[512];  
/* ... */  
rc = recv(connfd, buf, 512, 0);
```

```
int rc, connfd;  
char buf[512] __attribute__((flexos_share));  
/* ... */  
rc = flexos_gate(liblwip, recv, connfd, buf, 512, 0);
```



Porting

FlexOS: an Operating System for Flexible Isolation

Enabling Flexible Builds: Porting API

```
int rc, connfd;  
char buf[512];  
/* ... */  
rc = recv(connfd, buf, 512, 0);
```

```
int rc, connfd;  
char buf[512] __attribute__((flexos_share));  
/* ... */  
rc = flexos_gate(liblwip, recv, connfd, buf, 512, 0);
```

```
int rc, connfd;  
char *buf[512] = shared_malloc(512);  
/* ... */  
rc = mpk_gate(0, 1, recv, connfd, buf, 512, 0);
```

Replace with MPK gate

Replace with shared heap allocation

Porting

Automatic gate
instantiation at
build time



FlexOS: an Operating System for Flexible Isolation

Enabling Flexible Builds: Porting API

```
int rc, connfd;  
char buf[512];  
/* ... */  
rc = recv(connfd, buf, 512, 0);
```

```
int rc, connfd;  
char buf[512] __attribute__((flexos_share));  
/* ... */  
rc = flexos_gate(liblwip, recv, connfd, buf, 512, 0);
```

```
int rc, connfd;  
char buf[512];  
/* ... */  
rc = recv(connfd, buf, 512, 0);
```

lwip + app

Replace with standard stack allocation and function call

Porting

Automatic gate
instantiation at
build time

Coccinelle



FlexOS: an Operating System for Flexible Isolation

A Few Numbers

Redis throughput on a total of 80 configuration, isolation with MPK, varying


- Number of compartments (1, 2, 3)
- The distribution of software components in compartments
- Software hardening (per compartment, on/off)

FlexOS: an Operating System for Flexible Isolation

A Few Numbers

Redis throughput on a total of 80 configuration, isolation with MPK, varying

- Number of compartments (1, 2, 3)
- The distribution of software components in compartments
- Software hardening (per compartment, on/off)



FlexOS: an Operating System for Flexible Isolation

Design Space Exploration

- How to explore the wide design space FlexOS gives access to?
- Security is hard to quantify
- Propose a **semi-automated exploration strategy** based on partially ordered sets

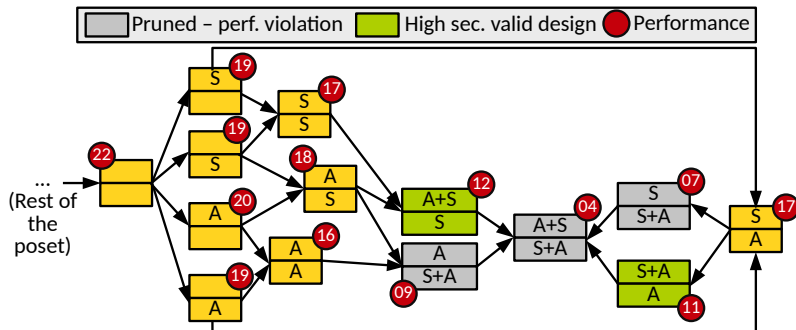
FlexOS: an Operating System for Flexible Isolation

Design Space Exploration

- How to explore the wide design space FlexOS gives access to?
- Security is hard to quantify
- Propose a **semi-automated exploration strategy** based on partially ordered sets

Example with 2 compartments and 2 software hardening techniques:

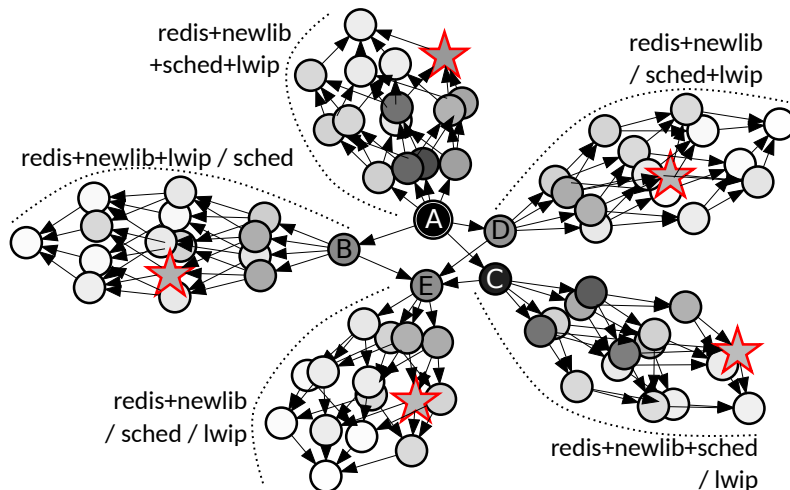
safe stack (S) and ASan (A):



FlexOS: an Operating System for Flexible Isolation

Design Space Exploration

Poset design space exploration for Redis:



Conclusion

Availability and Publications

Everything is open source!

- **HermiTux**

- <https://ssrg-vt.github.io/hermitux/>
- P. Olivier, D. Chiba, S. Lankes, C. Min and B. Ravindran, **A Binary-Compatible Unikernel**, *VEE'19*
- P. Olivier, H. Lefeuvre, D. Chiba, S. Lankes, C. Min and B. Ravindran, **A Syscall-Level Binary-Compatible Unikernel**, *IEEE TC*, 2021

- **FlexOS**

- <https://project-flexos.github.io/>
- H. Lefeuvre, V. Badoiu, A. Jung, S Teodorescu, S. Rauch, F. Huici, C. Raiciu, and P. Olivier, **FlexOS: Towards Flexible OS Isolation**, *ASPLOS'22*
- H. Lefeuvre, V. Badoiu, S Teodorescu, P. Olivier, T. Mosnoi, R. Deaconescu, F. Huici, and C. Raiciu, **FlexOS: Making OS Isolation Flexible**, *HotOS'21*
- **FlexOS is part of Hugo Lefeuvre's PhD work:** <https://www.owl.eu.com/research.html>

Please do not hesitate to get in touch :) pierre.olivier@manchester.ac.uk